Pandoc User's Guide

John MacFarlane

August 30, 2014

Chapter 1

Synopsis

pandoc [*options*] [*input-file*]...

Chapter 2

Description

Pandoc is a Haskell library for converting from one markup format to another, and a command-line tool that uses this library. It can read markdown and (subsets of) Textile, reStructuredText, HTML, LaTeX, MediaWiki markup, TWiki markup, Haddock markup, OPML, Emacs Org-mode, DocBook, txt2tags, EPUB and Word docx; and it can write plain text, markdown, reStructuredText, XHTML, HTML 5, LaTeX (including beamer slide shows), ConTeXt, RTF, OPML, DocBook, OpenDocument, ODT, Word docx, GNU Texinfo, MediaWiki markup, DokuWiki markup, Haddock markup, EPUB (v2 or v3), FictionBook2, Textile, groff man pages, Emacs Org-Mode, AsciiDoc, InDesign ICML, and Slidy, Slideous, DZSlides, reveal.js or S5 HTML slide shows. It can also produce PDF output on systems where LaTeX is installed.

Pandoc's enhanced version of markdown includes syntax for footnotes, tables, flexible ordered lists, definition lists, fenced code blocks, superscript, subscript, strikeout, title blocks, automatic tables of contents, embedded LaTeX math, citations, and markdown inside HTML block elements. (These enhancements, described below under Pandoc's markdown, can be disabled using the markdown_strict input or output format.)

In contrast to most existing tools for converting markdown to HTML, which use regex substitutions, Pandoc has a modular design: it consists of a set of readers, which parse text in a given format and produce a native representation of the document, and a set of writers, which convert this native representation into a target format. Thus, adding an input or output format requires only adding a reader or writer.

Using pandoc

If no *input-file* is specified, input is read from *stdin*. Otherwise, the *input-files* are concatenated (with a blank line between each) and used as input. Output goes to *stdout* by default (though output to *stdout* is disabled for the odt, docx, epub, and epub3 output formats). For output to a file, use the −o option:

```
pandoc −o output.html input.txt
```

By default, pandoc produces a document fragment, not a standalone document with a proper header and footer. To produce a standalone document, use the −s or −−standalone flag:

```
pandoc −s −o output.html input.txt
```

For more information on how standalone documents are produced, see Templates, below.

Instead of a file, an absolute URI may be given. In this case pandoc will fetch the content using HTTP:

```
pandoc −f html −t markdown http://www.fsf.org
```

If multiple input files are given, pandoc will concatenate them all (with blank lines between them) before parsing. This feature is disabled for binary input formats such as EPUB and docx.

The format of the input and output can be specified explicitly using command-line options. The input format can be specified using the −r/−−read or −f/−−from options, the output format using the −w/−−write or −t/−−to options. Thus, to convert hello.txt from markdown to LaTeX, you could type:

```
pandoc −f markdown −t latex hello.txt
```

To convert hello.html from html to markdown:

```
pandoc −f html −t markdown hello.html
```

Supported output formats are listed below under the −t/−−to option. Supported input formats are listed below under the −f/−−from option. Note that the rst, textile, latex, and html readers are not complete; there are some constructs that they do not parse.

If the input or output format is not specified explicitly, pandoc will attempt to guess it from the extensions of the input and output filenames. Thus, for example,

```
pandoc −o hello.tex hello.txt
```

will convert hello.txt from markdown to LaTeX. If no output file is specified (so that output goes to *stdout*), or if the output file's extension is unknown, the output format will default to HTML. If no input file is specified (so that input comes from *stdin*), or if the input files' extensions are unknown, the input format will be assumed to be markdown unless explicitly specified.

Pandoc uses the UTF-8 character encoding for both input and output. If your local character encoding is not UTF-8, you should pipe input and output through iconv:

```
iconv −t utf−8 input.txt | pandoc | iconv −f utf−8
```

Note that in some output formats (such as HTML, LaTeX, ConTeXt, RTF, OPML, DocBook, and Texinfo), information about the character encoding is included in the document header, which will only be included if you use the −s/−−standalone option.

Creating a PDF

Earlier versions of pandoc came with a program, markdown2pdf, that used pandoc and pdflatex to produce a PDF. This is no longer needed, since pandoc can now produce pdf output itself. To produce a PDF, simply specify an output file with a .pdf extension. Pandoc will create a latex file and use pdflatex (or another engine, see −−latex−engine) to convert it to PDF:

```
pandoc test.txt −o test.pdf
```

Production of a PDF requires that a LaTeX engine be installed (see −−latex−engine, below), and assumes that the following LaTeX packages are available: amssymb, amsmath, ifxetex, ifluatex, listings (if the −−listings option is used), fancyvrb, longtable, booktabs, url, graphicx, hyperref, ulem, babel (if the lang variable is set), fontspec (if xelatex or lualatex is used as the LaTeX engine), xltxtra and xunicode (if xelatex is used).

hsmarkdown

A user who wants a drop-in replacement for Markdown.pl may create a symbolic link to the pandoc executable called hsmarkdown. When invoked under the name hsmarkdown, pandoc will behave as if invoked with −f markdown_strict −−email−obfuscation=references, and all command-line options will be treated as regular arguments. However, this approach does not work under Cygwin, due to problems with its simulation of symbolic links.

Chapter 3

Options

General options

Specify input format. *FORMAT* can be native (native Haskell), json (JSON version of native AST), markdown (pandoc's extended markdown), markdown_strict (original unextended markdown), markdown_phpextra (PHP Markdown Extra extended markdown), markdown_github (github extended markdown), textile (Textile), rst (reStructuredText), html (HTML), docbook (DocBook), t2t (txt2tags), docx (docx), epub (EPUB), opml (OPML), org (Emacs Org-mode), mediawiki (MediaWiki markup), twiki (TWiki markup), haddock (Haddock markup), or latex (LaTeX). If +lhs is appended to markdown, rst, latex, or html, the input will be treated as literate Haskell source: see Literate Haskell support, below. Markdown syntax extensions can be individually enabled or disabled by appending +EXTENSION or −EXTENSION to the format name. So, for example, markdown_strict+footnotes+definition_lists is strict markdown with footnotes and definition lists enabled, and markdown−pipe_tables+hard_line_breaks is pandoc's markdown without pipe tables and with hard line breaks. See Pandoc's markdown, below, for a list of extensions and their names.

Specify output format. *FORMAT* can be native (native Haskell), json (JSON version of native AST), plain (plain text), markdown (pandoc's extended markdown), markdown_strict (original unextended markdown), markdown_phpextra (PHP Markdown extra extended markdown), markdown_github (github extended markdown), rst (reStructuredText), html (XHTML 1), html5 (HTML 5), latex (LaTeX), beamer (LaTeX beamer slide show), context (ConTeXt), man (groff man), mediawiki (MediaWiki markup), dokuwiki (DokuWiki markup), textile (Textile), org (Emacs Org-Mode), texinfo (GNU Texinfo), opml (OPML), docbook (DocBook), opendocument (OpenDocument), odt (OpenOffice text document), docx (Word docx), haddock (Haddock markup), rtf (rich text format), epub (EPUB v2 book), epub3 (EPUB v3), fb2 (Fiction-Book2 e-book), asciidoc (AsciiDoc), icml (InDesign ICML), slidy (Slidy HTML and javascript slide show), slideous (Slideous HTML and javascript slide show), dzslides (DZSlides HTML5 + javascript slide show), revealjs (reveal.js HTML5 + javascript slide show), s5 (S5 HTML and javascript slide show), or the path of a custom lua writer (see Custom writers, below). Note that odt, epub, and epub3 output will not be directed to *stdout*; an output filename must be specified using the −o/−−output option. If +lhs is appended to markdown, rst, latex, beamer, html, or html5, the output will be rendered as literate Haskell source: see Literate Haskell support, below. Markdown syntax extensions can be individually enabled or disabled by appending +EXTENSION or −EXTENSION to the format name, as described above under −f.

Write output to *FILE* instead of *stdout*. If *FILE* is −, output will go to *stdout*. (Exception: if the output format is odt, docx, epub, or epub3, output to stdout is disabled.)

Specify the user data directory to search for pandoc data files. If this option is not specified, the default user data directory will be used. This is

$HOME/.pandoc

in unix,

C:\ Documents And Settings\USERNAME\ Application Data\pandoc

in Windows XP, and

C:\ Users\USERNAME\AppData\Roaming\pandoc

in Windows 7. (You can find the default user data directory on your system by looking at the output of pandoc −−version.) A reference.odt, reference.docx, default.csl, epub.css, templates, slidy, slideous, or s5 directory placed in this directory will override pandoc's normal defaults.

Print version.

Show usage message.

Reader options

Parse untranslatable HTML codes and LaTeX environments as raw HTML or LaTeX, instead of ignoring them. Affects only HTML and LaTeX input. Raw HTML can be printed in markdown, reStructuredText, HTML, Slidy, Slideous, DZSlides, reveal.js, and S5 output; raw LaTeX can be printed in markdown, reStructuredText, LaTeX, and ConTeXt output. The default is for the readers to omit untranslatable HTML codes and LaTeX environments. (The LaTeX reader does pass through untranslatable LaTeX *commands*, even if −R is not specified.)

Produce typographically correct output, converting straight quotes to curly quotes, −−− to em-dashes, −− to en-dashes, and ... to ellipses. Nonbreaking spaces are inserted after certain abbreviations, such as "Mr." (Note: This option is significant only when the input format is markdown, markdown_strict, textile or twiki. It is selected automatically when the input format is textile or the output format is latex or context, unless −−no−tex−ligatures is used.)

Selects the pandoc <= 1.8.2.1 behavior for parsing smart dashes: − before a numeral is an en-dash, and −− is an em-dash. This option is selected automatically for textile input.

Specify the base level for headers (defaults to 1).

Specify classes to use for indented code blocks–for example, perl,numberLines or haskell. Multiple classes may be separated by spaces or commas.

Specify a default extension to use when image paths/URLs have no extension. This allows you to use the same source for formats that require different kinds of images. Currently this option only affects the markdown and LaTeX readers.

Specify an executable to be used as a filter transforming the Pandoc AST after the input is parsed and before the output is written. The executable should read JSON from stdin and write JSON to stdout. The JSON must be formatted like pandoc's own JSON input and output. The name of the output format will be passed to the filter as the first argument. Hence,

pandoc −−filter ./caps.py −t latex

is equivalent to

pandoc −t json | ./caps.py latex | pandoc −f json −t latex

The latter form may be useful for debugging filters.

Filters may be written in any language. Text.Pandoc.JSON exports toJSONFilter to facilitate writing filters in Haskell. Those who would prefer to write filters in python can use the module pandocfilters, installable from PyPI. See http://github.com/jgm/pandocfilters for the module and several examples. Note that the *EXECUTABLE* will be sought in the user's PATH, and not in the working directory, if no directory is provided. If you want to run a script in the working directory, preface the filename with ./.

Set the metadata field *KEY* to the value *VAL*. A value specified on the command line overrides a value specified in the document. Values will be parsed as YAML boolean or string values. If no value is specified, the value will be treated as Boolean true. Like −−variable, −−metadata causes template variables to be set. But unlike −−variable, −−metadata affects the metadata of the underlying document (which is accessible from filters and may be printed in some output formats).

Normalize the document after reading: merge adjacent Str or Emph elements, for example, and remove repeated Spaces.

Preserve tabs instead of converting them to spaces (the default). Note that this will only affect tabs in literal code spans and code blocks; tabs in regular text will be treated as spaces.

Specify the number of spaces per tab (default is 4).

Specifies what to do with insertions and deletions produced by the MS Word "track-changes" feature. *accept* (the default), inserts all insertions, and ignores all deletions. *reject* inserts all deletions and ignores insertions. *all* puts in both insertions and deletions, wrapped in spans with insertion and deletion classes, respectively. The author and time of change is included. *all* is useful for scripting: only accepting changes from a certain reviewer, say, or before a certain date. This option only affects the docx reader.

Extract images and other media contained in a docx or epub container to the path *DIR*, creating it if necessary, and adjust the images references in the document so they point to the extracted files. This option only affects the docx and epub readers.

General writer options

Produce output with an appropriate header and footer (e.g. a standalone HTML, LaTeX, or RTF file, not a fragment). This option is set automatically for pdf, epub, epub3, fb2, docx, and odt output.

Use *FILE* as a custom template for the generated document. Implies −−standalone. See Templates below for a description of template syntax. If no extension is specified, an extension corresponding to the writer will be added, so that −−template=special looks for special.html for HTML output. If the template is not found, pandoc will search for it in the user data directory (see −−data−dir). If this option is not used, a default template appropriate for the output format will be used (see −D/−−print−default−template).

Set the template variable *KEY* to the value *VAL* when rendering the document in standalone mode. This is generally only useful when the −−template option is used to specify a custom template, since pandoc automatically sets the variables used in the default templates. If no *VAL* is specified, the key will be given the value true.

Print the default template for an output *FORMAT*. (See −t for a list of possible *FORMAT*s.)

Print a default data file.

Disable text wrapping in output. By default, text is wrapped appropriately for the output format.

Specify length of lines in characters (for text wrapping).

Include an automatically generated table of contents (or, in the case of latex, context, and rst, an instruction to create one) in the output document. This option has no effect on man, docbook, slidy, slideous, s5, docx, or odt output.

Specify the number of section levels to include in the table of contents. The default is 3 (which means that level 1, 2, and 3 headers will be listed in the contents).

Disables syntax highlighting for code blocks and inlines, even when a language attribute is given.

Specifies the coloring style to be used in highlighted source code. Options are pygments (the default), kate, monochrome, espresso, zenburn, haddock, and tango.

Include contents of *FILE*, verbatim, at the end of the header. This can be used, for example, to include special CSS or javascript in HTML documents. This option can be used repeatedly to include multiple files in the header. They will be included in the order specified. Implies −−standalone.

Include contents of *FILE*, verbatim, at the beginning of the document body (e.g. after the <body> tag in HTML, or the \begin{document} command in LaTeX). This can be used to include navigation bars or banners in HTML documents. This option can be used repeatedly to include multiple files. They will be included in the order specified. Implies −−standalone.

Include contents of *FILE*, verbatim, at the end of the document body (before the </body> tag in HTML, or the \end{document} command in LaTeX). This option can be be used repeatedly to include multiple files. They will be included in the order specified. Implies −−standalone.

Options affecting specific writers

Produce a standalone HTML file with no external dependencies, using data: URIs to incorporate the contents of linked scripts, stylesheets, images, and videos. The resulting file should be "self-contained," in the sense that it needs no external files and no net access to be displayed properly by a browser. This option works only with HTML output formats, including html, html5, html+lhs, html5+lhs, s5, slidy, slideous, dzslides, and revealjs. Scripts, images, and stylesheets at absolute URLs will be downloaded; those at relative URLs will be sought relative to the working directory (if the first source file is local) or relative to the base URL (if the first source file is remote). −−self−contained does not work with −−mathjax.

Deprecated synonym for −−self−contained.

Produce HTML5 instead of HTML4. This option has no effect for writers other than html. (*Deprecated:* Use the html5 output format instead.)

Use <q> tags for quotes in HTML.

Use only ascii characters in output. Currently supported only for HTML output (which uses numerical entities instead of UTF-8 when this option is selected).

Use reference-style links, rather than inline links, in writing markdown or reStructuredText. By default inline links are used.

Use ATX style headers in markdown and asciidoc output. The default is to use setext-style headers for levels 1-2, and then ATX headers.

Treat top-level headers as chapters in LaTeX, ConTeXt, and DocBook output. When the LaTeX template uses the report, book, or memoir class, this option is implied. If beamer is the output format, top-level headers will become \part{..}.

Number section headings in LaTeX, ConTeXt, HTML, or EPUB output. By default, sections are not numbered. Sections with class unnumbered will never be numbered, even if −−number−sections is specified.

Offset for section headings in HTML output (ignored in other output formats). The first number is added to the section number for top-level headers, the second for second-level headers, and so on. So, for example, if you want the first top-level header in your document to be numbered "6", specify −−number−offset=5. If your document starts with a level-2 header which you want to be numbered "1.5", specify −−number−offset=1,4. Offsets are 0 by default. Implies −−number−sections.

Do not convert quotation marks, apostrophes, and dashes to the TeX ligatures when writing LaTeX or ConTeXt. Instead, just use literal unicode characters. This is needed for using advanced OpenType features with XeLaTeX and LuaLaTeX. Note: normally −−smart is selected automatically for LaTeX and ConTeXt output, but it must be specified explicitly if −−no−tex−ligatures is selected. If you use literal curly quotes, dashes, and ellipses in your source, then you may want to use −−no−tex−ligatures without −−smart.

Use listings package for LaTeX code blocks

Make list items in slide shows display incrementally (one by one). The default is for lists to be displayed all at once.

Specifies that headers with the specified level create slides (for beamer, s5, slidy, slideous, dzslides). Headers above this level in the hierarchy are used to divide the slide show into sections; headers below this level create subheads within a slide. The default is to set the slide level based on the contents of the document; see Structuring the slide show, below.

Wrap sections in <div> tags (or <section> tags in HTML5), and attach identifiers to the enclosing <div> (or <section>) rather than the header itself. See Section identifiers, below.

Specify a method for obfuscating mailto: links in HTML documents. *none* leaves mailto: links as they are. *javascript* obfuscates them using javascript. *references* obfuscates them by printing their letters as decimal or hexadecimal character references.

Specify a prefix to be added to all automatically generated identifiers in HTML and DocBook output, and to footnote numbers in markdown output. This is useful for preventing duplicate identifiers when generating fragments to be included in other pages.

Specify *STRING* as a prefix at the beginning of the title that appears in the HTML header (but not in the title as it appears at the beginning of the HTML body). Implies −−standalone.

Link to a CSS style sheet. This option can be be used repeatedly to include multiple files. They will be included in the order specified.

Use the specified file as a style reference in producing an ODT. For best results, the reference ODT should be a modified version of an ODT produced using pandoc. The contents of the reference ODT are ignored, but its stylesheets are used in the new ODT. If no reference ODT is specified on the command line, pandoc will look for a file reference.odt in the user data directory (see −−data−dir). If this is not found either, sensible defaults will be used.

Use the specified file as a style reference in producing a docx file. For best results, the reference docx should be a modified version of a docx file produced using pandoc. The contents of the reference docx are ignored, but its stylesheets and document properties (including margins, page size, header, and footer) are used in the new docx. If no reference docx is specified on the command line, pandoc will look for a file reference.docx in the user data directory (see −−data−dir). If this is not found either, sensible defaults will be used. The following styles are used by pandoc: [paragraph] Normal, Compact, Title, Subtitle, Authors, Date, Abstract, Heading 1, Heading 2, Heading 3, Heading 4, Heading 5, Block Quote, Definition Term, Definition, Bibliography, Body Text, Table Caption, Image Caption; [character] Default Paragraph Font, Body Text Char, Verbatim Char, Footnote Ref, Link.

Use the specified CSS file to style the EPUB. If no stylesheet is specified, pandoc will look for a file epub.css in the user data directory (see −−data−dir). If it is not found there, sensible defaults will be used.

Use the specified image as the EPUB cover. It is recommended that the image be less than 1000px in width and height. Note that in a markdown source document you can also specify cover−image in a YAML metadata block (see EPUB Metadata, below).

Look in the specified XML file for metadata for the EPUB. The file should contain a series of Dublin Core elements, as documented at http://dublincore.org/documents/dces/. For example:

```
<dc:rights>Creative Commons</dc:rights>
<dc:language>es−AR</dc:language>
```

By default, pandoc will include the following metadata elements: <dc:title> (from the document title), <dc:creator> (from the document authors), <dc:date> (from the document date, which should be in ISO 8601 format), <dc:language> (from the lang variable, or, if is not set, the locale), and <dc:identifier id="BookId"> (a randomly generated UUID). Any of these may be overridden by elements in the metadata file.

Note: if the source document is markdown, a YAML metadata block in the document can be used instead. See below under EPUB Metadata.

Embed the specified font in the EPUB. This option can be repeated to embed multiple fonts. To use embedded fonts, you will need to add declarations like the following to your CSS (see −−epub−stylesheet):

```
@font−face {
font−family: DejaVuSans;
font−style: normal;
font−weight: normal;
```

```
src:url("DejaVuSans−Regular.ttf");
}
@font−face {
font−family: DejaVuSans;
font−style: normal;
font−weight: bold;
src:url("DejaVuSans−Bold.ttf");
}
@font−face {
font−family: DejaVuSans;
font−style: italic;
font−weight: normal;
src:url("DejaVuSans−Oblique.ttf");
}
@font−face {
font−family: DejaVuSans;
font−style: italic;
font−weight: bold;
src:url("DejaVuSans−BoldOblique.ttf");
}
body { font−family: "DejaVuSans"; }
```

Specify the heading level at which to split the EPUB into separate "chapter" files. The default is to split into chapters at level 1 headers. This option only affects the internal composition of the EPUB, not the way chapters and sections are displayed to users. Some readers may be slow if the chapter files are too large, so for large documents with few level 1 headers, one might want to use a chapter level of 2 or 3.

Use the specified LaTeX engine when producing PDF output. The default is pdflatex. If the engine is not in your PATH, the full path of the engine may be specified here.

Citation rendering

Set the bibliography field in the document's metadata to *FILE*, overriding any value set in the metadata, and process citations using pandoc−citeproc. (This is equivalent to −−metadata bibliography=FILE −−filter pandoc−citeproc.)
Set the csl field in the document's metadata to *FILE*, overriding any value set in the metadata. (This is equivalent to −−metadata csl=FILE.)
Set the citation−abbreviations field in the document's metadata to *FILE*, overriding any value set in the metadata. (This is equivalent to −−metadata citation−abbreviations=FILE.)
Use natbib for citations in LaTeX output. This option is not for use with the pandoc−citeproc filter or with PDF output. It is intended for use in producing a LaTeX file that can be processed with pdflatex and bibtex.
Use biblatex for citations in LaTeX output. This option is not for use with the pandoc−citeproc filter or with PDF output. It is intended for use in producing a LaTeX file that can be processed with pdflatex and bibtex or biber.

Math rendering in HTML

Use the [LaTeXMathML](#) script to display embedded TeX math in HTML output. To insert a link to a local copy of the LaTeXMathML.js script, provide a *URL*. If no *URL* is provided, the contents of the script will be inserted directly into the HTML header, preserving portability at the price of efficiency. If you plan to use math on several pages, it is much better to link to a copy of the script, so it can be cached.
Convert TeX math to MathML (in docbook as well as html and html5). In standalone html output, a small javascript (or a link to such a script if a *URL* is supplied) will be inserted that allows the MathML to be viewed on some browsers.

Use jsMath to display embedded TeX math in HTML output. The *URL* should point to the jsMath load script (e.g. jsMath/easy/load.js); if provided, it will be linked to in the header of standalone HTML documents. If a *URL* is not provided, no link to the jsMath load script will be inserted; it is then up to the author to provide such a link in the HTML template.

Use MathJax to display embedded TeX math in HTML output. The *URL* should point to the MathJax.js load script. If a *URL* is not provided, a link to the MathJax CDN will be inserted.

Enclose TeX math in <eq> tags in HTML output. These can then be processed by gladTeX to produce links to images of the typeset formulas.

Render TeX math using the mimeTeX CGI script. If *URL* is not specified, it is assumed that the script is at /cgi−bin/mimetex.cgi.

Render TeX formulas using an external script that converts TeX formulas to images. The formula will be concatenated with the URL provided. If *URL* is not specified, the Google Chart API will be used.

−−katex[=*URL*] : Use KaTeX to display embedded TeX math in HTML output. The *URL* should point to the katex.js load script. If a *URL* is not provided, a link to the KaTeX CDN will be inserted.

−−katex−stylesheet=∗URL∗ : The *URL* should point to the katex.css stylesheet. If this option is not specified, a link to the KaTeX CDN will be inserted. Note that this option does not imply −−katex.

Options for wrapper scripts

Print information about command-line arguments to *stdout*, then exit. This option is intended primarily for use in wrapper scripts. The first line of output contains the name of the output file specified with the −o option, or − (for *stdout*) if no output file was specified. The remaining lines contain the command-line arguments, one per line, in the order they appear. These do not include regular Pandoc options and their arguments, but do include any options appearing after a −− separator at the end of the line.

Ignore command-line arguments (for use in wrapper scripts). Regular Pandoc options are not ignored. Thus, for example,

```
pandoc −−ignore−args −o foo.html −s foo.txt −− −e latin1
```

is equivalent to

```
pandoc −o foo.html −s
```

Chapter 4

Templates

When the −s/−−standalone option is used, pandoc uses a template to add header and footer material that is needed for a self-standing document. To see the default template that is used, just type

`pandoc −D FORMAT`

where FORMAT is the name of the output format. A custom template can be specified using the −−template option. You can also override the system default templates for a given output format FORMAT by putting a file templates/default.FORMAT in the user data directory (see −−data−dir, above). *Exceptions:* For odt output, customize the default.opendocument template. For pdf output, customize the default.latex template.

Templates may contain *variables*. Variable names are sequences of alphanumerics, −, and _, starting with a letter. A variable name surrounded by $ signs will be replaced by its value. For example, the string $title$ in

`<title>$title$</title>`

will be replaced by the document title.

To write a literal $ in a template, use $$.

Some variables are set automatically by pandoc. These vary somewhat depending on the output format, but include metadata fields (such as title, author, and date) as well as the following:

contents specified by −H/−−include−in−header (may have multiple values)
non-null value if −−toc/−−table−of−contents was specified
contents specified by −B/−−include−before−body (may have multiple values)
contents specified by −A/−−include−after−body (may have multiple values)
body of document
language code for HTML or LaTeX documents
base URL for Slidy documents (defaults to http://www.w3.org/Talks/Tools/Slidy2)
base URL for Slideous documents (defaults to slideous)
base URL for S5 documents (defaults to s5/default)
base URL for reveal.js documents (defaults to reveal.js)
reveal.js or LaTeX beamer theme
reveal.js transition
font size (10pt, 11pt, 12pt) for LaTeX documents
document class for LaTeX documents
option for LaTeX documentclass, e.g. oneside; may be repeated for multiple options
options for LaTeX geometry class, e.g. margin=1in; may be repeated for multiple options
adjusts line spacing (requires the setspace package)
font package to use for LaTeX documents (with pdflatex): TeXLive has bookman (Bookman), utopia or fourier (Utopia), fouriernc (New Century Schoolbook), times or txfonts (Times), mathpazo or pxfonts or mathpple (Palatino), libertine (Linux Libertine), arev (Arev Sans), and the default lmodern, among others.

> fonts for LaTeX documents (works only with xelatex and lualatex)
> colortheme for LaTeX beamer documents
> fonttheme for LaTeX beamer documents
> color for internal links in LaTeX documents (red, green, magenta, cyan, blue, black)
> color for external links in LaTeX documents
> color for citation links in LaTeX documents
> causes links to be printed as footnotes in LaTeX documents
> include table of contents in LaTeX documents
> level of section to include in table of contents in LaTeX documents
> include list of figures in LaTeX documents
> include list of tables in LaTeX documents
> bibliography style in LaTeX, when used with −−natbib
> bibliography files to use in LaTeX, with −−natbib or −−biblatex
> section number in man pages
> header in man pages
> footer in man pages

Variables may be set at the command line using the −V/−−variable option. Variables set in this way override metadata fields with the same name.

Templates may contain conditionals. The syntax is as follows:

```
$if(variable)$
X
$else$
Y
$endif$
```

This will include X in the template if variable has a non-null value; otherwise it will include Y. X and Y are placeholders for any valid template text, and may include interpolated variables or other conditionals. The $else$ section may be omitted.

When variables can have multiple values (for example, author in a multi-author document), you can use the for keyword:

```
$for(author)$
<meta name="author" content="$author$" />
$endfor$
```

You can optionally specify a separator to be used between consecutive items:

```
$for(author)$$author$$sep$, $endfor$
```

A dot can be used to select a field of a variable that takes an object as its value. So, for example:

```
$author.name$ ($author.affiliation$)
```

If you use custom templates, you may need to revise them as pandoc changes. We recommend tracking the changes in the default templates, and modifying your custom templates accordingly. An easy way to do this is to fork the pandoc-templates repository (http://github.com/jgm/pandoc-templates) and merge in changes after each pandoc release.

Chapter 5

Pandoc's markdown

Pandoc understands an extended and slightly revised version of John Gruber's markdown syntax. This document explains the syntax, noting differences from standard markdown. Except where noted, these differences can be suppressed by using the markdown_strict format instead of markdown. An extensions can be enabled by adding +EXTENSION to the format name and disabled by adding −EXTENSION. For example, markdown_strict+footnotes is strict markdown with footnotes enabled, while markdown−footnotes−pipe_tables is pandoc's markdown without footnotes or pipe tables.

Philosophy

Markdown is designed to be easy to write, and, even more importantly, easy to read:

> ▓▓▓ is, as plain text, without looking like it's been marked up with tags or formatting instructions. – John Gruber

This principle has guided pandoc's decisions in finding syntax for tables, footnotes, and other extensions.

There is, however, one respect in which pandoc's aims are different from the original aims of markdown. Whereas markdown was originally designed with HTML generation in mind, pandoc is designed for multiple output formats. Thus, while pandoc allows the embedding of raw HTML, it discourages it, and provides other, non-HTMLish ways of representing important document elements like definition lists, tables, mathematics, and footnotes.

Paragraphs

A paragraph is one or more lines of text followed by one or more blank line. Newlines are treated as spaces, so you can reflow your paragraphs as you like. If you need a hard line break, put two or more spaces at the end of a line.

Extension: escaped_line_breaks

A backslash followed by a newline is also a hard line break. Note: in multiline and grid table cells, this is the only way to create a hard line break, since trailing spaces in the cells are ignored.

Headers

There are two kinds of headers, Setext and atx.

Setext-style headers

A setext-style header is a line of text "underlined" with a row of = signs (for a level one header) or − signs (for a level two header):

```
A level−one header
==================
```

```
A level−two header
−−−−−−−−−−−−−−−−−−−
```

The header text can contain inline formatting, such as emphasis (see Inline formatting, below).

Atx-style headers

An Atx-style header consists of one to six # signs and a line of text, optionally followed by any number of # signs. The number of # signs at the beginning of the line is the header level:

```
## A level−two header
```

```
### A level−three header ###
```

As with setext-style headers, the header text can contain formatting:

```
# A level−one header with a [link](/url) and *emphasis*
```

Extension: blank_before_header

Standard markdown syntax does not require a blank line before a header. Pandoc does require this (except, of course, at the beginning of the document). The reason for the requirement is that it is all too easy for a # to end up at the beginning of a line by accident (perhaps through line wrapping). Consider, for example:

```
I like several of their flavors of ice cream:
#22, for example, and #5.
```

Header identifiers in HTML, LaTeX, and ConTeXt

Extension: header_attributes

Headers can be assigned attributes using this syntax at the end of the line containing the header text:

```
{#identifier .class .class key=value key=value}
```

Thus, for example, the following headers will all be assigned the identifier foo:

```
# My header {#foo}
```

```
## My header ##    {#foo}
```

```
My other header   {#foo}
−−−−−−−−−−−−−−−−−
```

(This syntax is compatible with PHP Markdown Extra.)

Note that although this syntax allows assignment of classes and key/value attributes, writers generally don't use all of this information. Identifiers, classes, and key/value attributes are used in HTML and HTML-based formats

such as EPUB and slidy. Identifiers are used for labels and link anchors in the LaTeX, ConTeXt, Textile, and AsciiDoc writers.

Headers with the class unnumbered will not be numbered, even if −−number−sections is specified. A single hyphen (−) in an attribute context is equivalent to .unnumbered, and preferable in non-English documents. So,

My header {−}

is just the same as

My header {.unnumbered}

Extension: auto_identifiers

A header without an explicitly specified identifier will be automatically assigned a unique identifier based on the header text. To derive the identifier from the header text,

- Remove all formatting, links, etc.
- Remove all footnotes.
- Remove all punctuation, except underscores, hyphens, and periods.
- Replace all spaces and newlines with hyphens.
- Convert all alphabetic characters to lowercase.
- Remove everything up to the first letter (identifiers may not begin with a number or punctuation mark).
- If nothing is left after this, use the identifier section.

Thus, for example,

Header	Identifier
Header identifiers in HTML	header−identifiers−in−html
Dogs?–in *my* house?	dogs−−in−my−house
HTML, S5, or RTF?	html−s5−or−rtf
3. Applications	applications
33	section

These rules should, in most cases, allow one to determine the identifier from the header text. The exception is when several headers have the same text; in this case, the first will get an identifier as described above; the second will get the same identifier with −1 appended; the third with −2; and so on.

These identifiers are used to provide link targets in the table of contents generated by the −−toc|−−table−of−contents option. They also make it easy to provide links from one section of a document to another. A link to this section, for example, might look like this:

See the section on
[header identifiers](#header−identifiers−in−html−latex−and−context).

Note, however, that this method of providing links to sections works only in HTML, LaTeX, and ConTeXt formats.

If the −−section−divs option is specified, then each section will be wrapped in a div (or a section, if −−html5 was specified), and the identifier will be attached to the enclosing <div> (or <section>) tag rather than the header itself. This allows entire sections to be manipulated using javascript or treated differently in CSS.

Extension: implicit_header_references

Pandoc behaves as if reference links have been defined for each header. So, instead of

[header identifiers](#header−identifiers−in−html)

you can simply write

```
[ header  identifiers ]
```

or

```
[ header  identifiers ] [ ]
```

or

```
[ the  section  on  header  identifiers ] [ header  identifiers ]
```

If there are multiple headers with identical text, the corresponding reference will link to the first one only, and you will need to use explicit links to link to the others, as described above.

Unlike regular reference links, these references are case-sensitive.

Note: if you have defined an explicit identifier for a header, then implicit references to it will not work.

Block quotations

Markdown uses email conventions for quoting blocks of text. A block quotation is one or more paragraphs or other block elements (such as lists or headers), with each line preceded by a > character and a space. (The > need not start at the left margin, but it should not be indented more than three spaces.)

```
> This  is  a  block  quote .  This
> paragraph  has  two  lines .
>
> 1.  This  is  a  list  inside  a  block  quote .
> 2.  Second  item .
```

A "lazy" form, which requires the > character only on the first line of each block, is also allowed:

```
> This  is  a  block  quote .  This
paragraph  has  two  lines .

> 1.  This  is  a  list  inside  a  block  quote .
2.  Second  item .
```

Among the block elements that can be contained in a block quote are other block quotes. That is, block quotes can be nested:

```
> This  is  a  block  quote .
>
> > A  block  quote  within  a  block  quote .
```

Extension: blank_before_blockquote

Standard markdown syntax does not require a blank line before a block quote. Pandoc does require this (except, of course, at the beginning of the document). The reason for the requirement is that it is all too easy for a > to end up at the beginning of a line by accident (perhaps through line wrapping). So, unless the markdown_strict format is used, the following does not produce a nested block quote in pandoc:

```
> This  is  a  block  quote .
>> Nested .
```

Verbatim (code) blocks

Indented code blocks

A block of text indented four spaces (or one tab) is treated as verbatim text: that is, special characters do not trigger special formatting, and all spaces and line breaks are preserved. For example,

```
if (a > 3) {
  moveShip(5 * gravity, DOWN);
}
```

The initial (four space or one tab) indentation is not considered part of the verbatim text, and is removed in the output.

Note: blank lines in the verbatim text need not begin with four spaces.

Fenced code blocks

Extension: fenced_code_blocks

In addition to standard indented code blocks, Pandoc supports *fenced* code blocks. These begin with a row of three or more tildes (~) or backticks (') and end with a row of tildes or backticks that must be at least as long as the starting row. Everything between these lines is treated as code. No indentation is necessary:

```
~~~~~~~
if (a > 3) {
  moveShip(5 * gravity, DOWN);
}
~~~~~~~
```

Like regular code blocks, fenced code blocks must be separated from surrounding text by blank lines.

If the code itself contains a row of tildes or backticks, just use a longer row of tildes or backticks at the start and end:

```
~~~~~~~~~~~~~~~~
~~~~~~~~~~
code including tildes
~~~~~~~~~~
~~~~~~~~~~~~~~~~
```

Extension: fenced_code_attributes

Optionally, you may attach attributes to the code block using this syntax:

```
~~~~ {#mycode .haskell .numberLines startFrom="100"}
qsort []     = []
qsort (x:xs) = qsort (filter (< x) xs) ++ [x] ++
               qsort (filter (>= x) xs)
~~~~~~~~~~~~~~~~~~~~~~~~~~~~~~~~~~~~~~~~~~~~~~~~~~~~
```

Here mycode is an identifier, haskell and numberLines are classes, and startFrom is an attribute with value 100. Some output formats can use this information to do syntax highlighting. Currently, the only output formats that uses this information are HTML and LaTeX. If highlighting is supported for your output format and language, then the code block above will appear highlighted, with numbered lines. (To see which languages are supported, do pandoc −−version.) Otherwise, the code block above will appear as follows:

```
<pre id="mycode" class="haskell numberLines" startFrom="100">
  <code>
  ...
  </code>
</pre>
```

A shortcut form can also be used for specifying the language of the code block:

```
```haskell
qsort [] = []
```
```

This is equivalent to:

```
``` {.haskell}
qsort [] = []
```
```

If the fenced_code_attributes extension is disabled, but input contains class attribute(s) for the codeblock, the first class attribute will be printed after the opening fence as a bare word.

To prevent all highlighting, use the $--$no$-$highlight flag. To set the highlighting style, use $--$highlight$-$style.

Line blocks

Extension: line_blocks

A line block is a sequence of lines beginning with a vertical bar ($|$) followed by a space. The division into lines will be preserved in the output, as will any leading spaces; otherwise, the lines will be formatted as markdown. This is useful for verse and addresses:

```
| The limerick packs laughs anatomical
| In space that is quite economical.
|    But the good ones I've seen
|    So seldom are clean
| And the clean ones so seldom are comical

| 200 Main St.
| Berkeley, CA 94718
```

The lines can be hard-wrapped if needed, but the continuation line must begin with a space.

```
| The Right Honorable Most Venerable and Righteous Samuel L.
  Constable, Jr.
| 200 Main St.
| Berkeley, CA 94718
```

This syntax is borrowed from reStructuredText.

Lists

Bullet lists

A bullet list is a list of bulleted list items. A bulleted list item begins with a bullet ($*$, $+$, or $-$). Here is a simple example:

```
* one
* two
* three
```

This will produce a "compact" list. If you want a "loose" list, in which each item is formatted as a paragraph, put spaces between the items:

```
* one

* two

* three
```

The bullets need not be flush with the left margin; they may be indented one, two, or three spaces. The bullet must be followed by whitespace.

List items look best if subsequent lines are flush with the first line (after the bullet):

```
* here is my first
  list item.
* and my second.
```

But markdown also allows a "lazy" format:

```
* here is my first
list item.
* and my second.
```

The four-space rule

A list item may contain multiple paragraphs and other block-level content. However, subsequent paragraphs must be preceded by a blank line and indented four spaces or a tab. The list will look better if the first paragraph is aligned with the rest:

```
  * First paragraph.

    Continued.

  * Second paragraph. With a code block, which must be indented
    eight spaces:

        { code }
```

List items may include other lists. In this case the preceding blank line is optional. The nested list must be indented four spaces or one tab:

```
* fruits
    + apples
        - macintosh
        - red delicious
    + pears
    + peaches
* vegetables
    + broccoli
    + chard
```

As noted above, markdown allows you to write list items "lazily," instead of indenting continuation lines. However, if there are multiple paragraphs or other blocks in a list item, the first line of each must be indented.

```
+ A lazy , lazy , list
item .

+ Another one ; this looks
bad but is legal .

    Second paragraph of second
list item .
```

Note: Although the four-space rule for continuation paragraphs comes from the official markdown syntax guide, the reference implementation, Markdown.pl, does not follow it. So pandoc will give different results than Markdown.pl when authors have indented continuation paragraphs fewer than four spaces.

The markdown syntax guide is not explicit whether the four-space rule applies to *all* block-level content in a list item; it only mentions paragraphs and code blocks. But it implies that the rule applies to all block-level content (including nested lists), and pandoc interprets it that way.

Ordered lists

Ordered lists work just like bulleted lists, except that the items begin with enumerators rather than bullets.

In standard markdown, enumerators are decimal numbers followed by a period and a space. The numbers themselves are ignored, so there is no difference between this list:

```
1.   one
2.   two
3.   three
```

and this one:

```
5.   one
7.   two
1.   three
```

Extension: fancy_lists

Unlike standard markdown, Pandoc allows ordered list items to be marked with uppercase and lowercase letters and roman numerals, in addition to arabic numerals. List markers may be enclosed in parentheses or followed by a single right-parentheses or period. They must be separated from the text that follows by at least one space, and, if the list marker is a capital letter with a period, by at least two spaces.[1]

The fancy_lists extension also allows '#' to be used as an ordered list marker in place of a numeral:

```
#.  one
#.  two
```

Extension: startnum

Pandoc also pays attention to the type of list marker used, and to the starting number, and both of these are preserved where possible in the output format. Thus, the following yields a list with numbers followed by a single parenthesis, starting with 9, and a sublist with lowercase roman numerals:

---

[1]The point of this rule is to ensure that normal paragraphs starting with people's initials, like B. Russell was an English philosopher. do not get treated as list items. This rule will not prevent (C) 2007 Joe Smith from being interpreted as a list item. In this case, a backslash escape can be used: (C) 2007 Joe Smith

```
 9)    Ninth
10)    Tenth
11)    Eleventh
         i.  subone
        ii.  subtwo
       iii.  subthree
```

Pandoc will start a new list each time a different type of list marker is used. So, the following will create three lists:

```
(2)  Two
(5)  Three
1.   Four
*    Five
```

If default list markers are desired, use #.:

```
#.   one
#.   two
#.   three
```

Definition lists

Extension: definition_lists

Pandoc supports definition lists, using the syntax of [PHP Markdown Extra](#) with some extensions.[2]

```
Term 1

:    Definition 1

Term 2 with *inline markup*

:    Definition 2

         { some code, part of Definition 2 }

     Third paragraph of definition 2.
```

Each term must fit on one line, which may optionally be followed by a blank line, and must be followed by one or more definitions. A definition begins with a colon or tilde, which may be indented one or two spaces.

A term may have multiple definitions, and each definition may consist of one or more block elements (paragraph, code block, list, etc.), each indented four spaces or one tab stop. The body of the definition (including the first line, aside from the colon or tilde) should be indented four spaces. However, as with other markdown lists, you can "lazily" omit indentation except at the beginning of a paragraph or other block element:

```
Term 1

:    Definition
with lazy continuation.

     Second paragraph of the definition.
```

If you leave space before the definition (as in the example above), the text of the definition will be treated as a paragraph. In some output formats, this will mean greater spacing between term/definition pairs. For a more compact definition list, omit the space before the definition:

---

[2]I have been influenced by the suggestions of [David Wheeler](#).

```
Term 1
  ~ Definition 1

Term 2
  ~ Definition 2a
  ~ Definition 2b
```

Note that space between items in a definition list is required. (A variant that loosens this requirement, but disallows "lazy" hard wrapping, can be activated with compact_definition_lists: see Non-pandoc extensions, below.)

Numbered example lists

Extension: example_lists

The special list marker @ can be used for sequentially numbered examples. The first list item with a @ marker will be numbered '1', the next '2', and so on, throughout the document. The numbered examples need not occur in a single list; each new list using @ will take up where the last stopped. So, for example:

```
(@)  My first example will be numbered (1).
(@)  My second example will be numbered (2).

Explanation of examples.

(@)  My third example will be numbered (3).
```

Numbered examples can be labeled and referred to elsewhere in the document:

```
(@good)  This is a good example.

As (@good) illustrates, ...
```

The label can be any string of alphanumeric characters, underscores, or hyphens.

Compact and loose lists

Pandoc behaves differently from Markdown.pl on some "edge cases" involving lists. Consider this source:

```
+   First
+   Second:
    −   Fee
    −   Fie
    −   Foe

+   Third
```

Pandoc transforms this into a "compact list" (with no <p> tags around "First", "Second", or "Third"), while markdown puts <p> tags around "Second" and "Third" (but not "First"), because of the blank space around "Third". Pandoc follows a simple rule: if the text is followed by a blank line, it is treated as a paragraph. Since "Second" is followed by a list, and not a blank line, it isn't treated as a paragraph. The fact that the list is followed by a blank line is irrelevant. (Note: Pandoc works this way even when the markdown_strict format is specified. This behavior is consistent with the official markdown syntax description, even though it is different from that of Markdown.pl.)

Ending a list

What if you want to put an indented code block after a list?

```
−     item  one
−     item  two

      {  my  code  block  }
```

Trouble! Here pandoc (like other markdown implementations) will treat { my code block } as the second paragraph of item two, and not as a code block.

To "cut off" the list after item two, you can insert some non-indented content, like an HTML comment, which won't produce visible output in any format:

```
−     item  one
−     item  two

<!−−  end  of  list  −−>

      {  my  code  block  }
```

You can use the same trick if you want two consecutive lists instead of one big list:

```
1.    one
2.    two
3.    three

<!−−  −−>

1.    uno
2.    dos
3.    tres
```

Horizontal rules

A line containing a row of three or more ∗, −, or _ characters (optionally separated by spaces) produces a horizontal rule:

```
*   *   *   *

_____
```

Tables

Four kinds of tables may be used. The first three kinds presuppose the use of a fixed-width font, such as Courier. The fourth kind can be used with proportionally spaced fonts, as it does not require lining up columns.

Extension: table_captions

A caption may optionally be provided with all 4 kinds of tables (as illustrated in the examples below). A caption is a paragraph beginning with the string Table: (or just :), which will be stripped off. It may appear either before or after the table.

Extension: simple_tables

Simple tables look like this:

Right	Left	Center	Default
12	12	12	12
123	123	123	123
1	1	1	1

Table: Demonstration of simple table syntax.

The headers and table rows must each fit on one line. Column alignments are determined by the position of the header text relative to the dashed line below it:[3]

- If the dashed line is flush with the header text on the right side but extends beyond it on the left, the column is right-aligned.
- If the dashed line is flush with the header text on the left side but extends beyond it on the right, the column is left-aligned.
- If the dashed line extends beyond the header text on both sides, the column is centered.
- If the dashed line is flush with the header text on both sides, the default alignment is used (in most cases, this will be left).

The table must end with a blank line, or a line of dashes followed by a blank line.

The column headers may be omitted, provided a dashed line is used to end the table. For example:

12	12	12	12
123	123	123	123
1	1	1	1

When headers are omitted, column alignments are determined on the basis of the first line of the table body. So, in the tables above, the columns would be right, left, center, and right aligned, respectively.

Extension: multiline_tables

Multiline tables allow headers and table rows to span multiple lines of text (but cells that span multiple columns or rows of the table are not supported). Here is an example:

Centered Header	Default Aligned	Right Aligned	Left Aligned
First	row	12.0	Example of a row that spans multiple lines.
Second	row	5.0	Here's another one. Note the blank line between rows.

Table: Here's the caption. It, too, may span multiple lines.

---
[3]This scheme is due to Michel Fortin, who proposed it on the Markdown discussion list.

These work like simple tables, but with the following differences:

- They must begin with a row of dashes, before the header text (unless the headers are omitted).
- They must end with a row of dashes, then a blank line.
- The rows must be separated by blank lines.

In multiline tables, the table parser pays attention to the widths of the columns, and the writers try to reproduce these relative widths in the output. So, if you find that one of the columns is too narrow in the output, try widening it in the markdown source.

Headers may be omitted in multiline tables as well as simple tables:

```
-------------  ----------  -------------------  ------------------------------
   First       row                       12.0   Example of a row that
                                                 spans multiple lines.

   Second      row                        5.0   Here's another one. Note
                                                 the blank line between
                                                 rows.
-------------  ----------  -------------------  ------------------------------

: Here's a multiline table without headers.
```

It is possible for a multiline table to have just one row, but the row should be followed by a blank line (and then the row of dashes that ends the table), or the table may be interpreted as a simple table.

Extension: grid_tables

Grid tables look like this:

```
: Sample grid table.

+---------------+---------------+--------------------+
| Fruit         | Price         | Advantages         |
+===============+===============+====================+
| Bananas       | $1.34         | - built-in wrapper |
|               |               | - bright color     |
+---------------+---------------+--------------------+
| Oranges       | $2.10         | - cures scurvy     |
|               |               | - tasty            |
+---------------+---------------+--------------------+
```

The row of =s separates the header from the table body, and can be omitted for a headerless table. The cells of grid tables may contain arbitrary block elements (multiple paragraphs, code blocks, lists, etc.). Alignments are not supported, nor are cells that span multiple columns or rows. Grid tables can be created easily using Emacs table mode.

Extension: pipe_tables

Pipe tables look like this:

```
| Right | Left | Default | Center |
|------:|:-----|---------|:------:|
|   12  | 12   |    12   |    12  |
|  123  | 123  |   123   |   123  |
|    1  | 1    |     1   |     1  |

  : Demonstration of pipe table syntax.
```

The syntax is [the same as in PHP markdown extra](). The beginning and ending pipe characters are optional, but pipes are required between all columns. The colons indicate column alignment as shown. The header can be omitted, but the horizontal line must still be included, as it defines column alignments.

Since the pipes indicate column boundaries, columns need not be vertically aligned, as they are in the above example. So, this is a perfectly legal (though ugly) pipe table:

```
fruit| price
-----|-----:
apple|2.05
pear|1.37
orange|3.09
```

The cells of pipe tables cannot contain block elements like paragraphs and lists, and cannot span multiple lines. Note also that in LaTeX/PDF output, the cells produced by pipe tables will not wrap, since there is no information available about relative widths. If you want content to wrap within cells, use multiline or grid tables.

Note: Pandoc also recognizes pipe tables of the following form, as can produced by Emacs' orgtbl-mode:

```
| One   | Two      |
|-----+--------|
| my    | table  |
| is    | nice   |
```

The difference is that + is used instead of |. Other orgtbl features are not supported. In particular, to get non-default column alignment, you'll need to add colons as above.

Metadata blocks

Extension: pandoc_title_block

If the file begins with a title block

```
% title
% author(s) (separated by semicolons)
% date
```

it will be parsed as bibliographic information, not regular text. (It will be used, for example, in the title of standalone LaTeX or HTML output.) The block may contain just a title, a title and an author, or all three elements. If you want to include an author but no title, or a title and a date but no author, you need a blank line:

```
%
% Author
```

```
% My title
%
% June 15, 2006
```

The title may occupy multiple lines, but continuation lines must begin with leading space, thus:

```
% My title
  on multiple lines
```

If a document has multiple authors, the authors may be put on separate lines with leading space, or separated by semicolons, or both. So, all of the following are equivalent:

```
% Author One
  Author Two
```

% Author One; Author Two

% Author One;
   Author Two

The date must fit on one line.

All three metadata fields may contain standard inline formatting (italics, links, footnotes, etc.).

Title blocks will always be parsed, but they will affect the output only when the −−standalone (−s) option is chosen. In HTML output, titles will appear twice: once in the document head – this is the title that will appear at the top of the window in a browser – and once at the beginning of the document body. The title in the document head can have an optional prefix attached (−−title−prefix or −T option). The title in the body appears as an H1 element with class "title", so it can be suppressed or reformatted with CSS. If a title prefix is specified with −T and no title block appears in the document, the title prefix will be used by itself as the HTML title.

The man page writer extracts a title, man page section number, and other header and footer information from the title line. The title is assumed to be the first word on the title line, which may optionally end with a (single-digit) section number in parentheses. (There should be no space between the title and the parentheses.) Anything after this is assumed to be additional footer and header text. A single pipe character (|) should be used to separate the footer text from the header text. Thus,

% PANDOC(1)

will yield a man page with the title PANDOC and section 1.

% PANDOC(1) Pandoc User Manuals

will also have "Pandoc User Manuals" in the footer.

% PANDOC(1) Pandoc User Manuals | Version 4.0

will also have "Version 4.0" in the header.

Extension: yaml_metadata_block

A YAML metadata block is a valid YAML object, delimited by a line of three hyphens (−−−) at the top and a line of three hyphens (−−−) or three dots (...) at the bottom. A YAML metadata block may occur anywhere in the document, but if it is not at the beginning, it must be preceded by a blank line. (Note that, because of the way pandoc concatenates input files when several are provided, you may also keep the metadata in a separate YAML file and pass it to pandoc as an argument, along with your markdown files:

pandoc chap1.md chap2.md chap3.md metadata.yaml −s −o book.html

Just be sure that the YAML file begins with −−− and ends with −−− or ... .)

Metadata will be taken from the fields of the YAML object and added to any existing document metadata. Metadata can contain lists and objects (nested arbitrarily), but all string scalars will be interpreted as markdown. Fields with names ending in an underscore will be ignored by pandoc. (They may be given a role by external processors.)

A document may contain multiple metadata blocks. The metadata fields will be combined through a *left-biased union*: if two metadata blocks attempt to set the same field, the value from the first block will be taken.

When pandoc is used with −t markdown to create a markdown document, a YAML metadata block will be produced only if the −s/−−standalone option is used. All of the metadata will appear in a single block at the beginning of the document.

Note that YAML escaping rules must be followed. Thus, for example, if a title contains a colon, it must be quoted. The pipe character (|) can be used to begin an indented block that will be interpreted literally, without need for escaping. This form is necessary when the field contains blank lines:

```
---
title:  'This is the title: it contains a colon'
author:
- name: Author One
  affiliation: University of Somewhere
- name: Author Two
  affiliation: University of Nowhere
tags: [nothing, nothingness]
abstract: |
  This is the abstract.

  It consists of two paragraphs.
...
```

Template variables will be set automatically from the metadata. Thus, for example, in writing HTML, the variable abstract will be set to the HTML equivalent of the markdown in the abstract field:

```
<p>This is the abstract.</p>
<p>It consists of two paragraphs.</p>
```

Note: The author variable in the default templates expects a simple list or string. To use the structured authors in the example, you would need a custom template. For example:

```
$for(author)$
$if(author.name)$
$author.name$$if(author.affiliation)$ ($author.affiliation$)$endif$
$else$
$author$
$endif$
$endfor$
```

Backslash escapes

Extension: all_symbols_escapable

Except inside a code block or inline code, any punctuation or space character preceded by a backslash will be treated literally, even if it would normally indicate formatting. Thus, for example, if one writes

```
*\*hello\**
```

one will get

```
<em>*hello*</em>
```

instead of

```
<strong>hello</strong>
```

This rule is easier to remember than standard markdown's rule, which allows only the following characters to be backslash-escaped:

```
\`*_{}[]()>#+-.!
```

(However, if the markdown_strict format is used, the standard markdown rule will be used.)

A backslash-escaped space is parsed as a nonbreaking space. It will appear in TeX output as ~ and in HTML and XML as \ or \ .

A backslash-escaped newline (i.e. a backslash occurring at the end of a line) is parsed as a hard line break. It will appear in TeX output as \\ and in HTML as <br />. This is a nice alternative to markdown's "invisible" way of indicating hard line breaks using two trailing spaces on a line.

Backslash escapes do not work in verbatim contexts.

Smart punctuation

Extension

If the −−smart option is specified, pandoc will produce typographically correct output, converting straight quotes to curly quotes, −−− to em-dashes, −− to en-dashes, and … to ellipses. Nonbreaking spaces are inserted after certain abbreviations, such as "Mr."

Note: if your LaTeX template uses the csquotes package, pandoc will detect automatically this and use \enquote{...} for quoted text.

Inline formatting

Emphasis

To *emphasize* some text, surround it with ∗s or _, like this:

```
This text is _emphasized with underscores_, and this
is *emphasized with asterisks*.
```

Double ∗ or _ produces **strong emphasis**:

```
This is **strong emphasis** and __with underscores__.
```

A ∗ or _ character surrounded by spaces, or backslash-escaped, will not trigger emphasis:

```
This is * not emphasized *, and \*neither is this\*.
```

Extension: intraword_underscores

Because _ is sometimes used inside words and identifiers, pandoc does not interpret a _ surrounded by alphanumeric characters as an emphasis marker. If you want to emphasize just part of a word, use ∗:

```
feas*ible*, not feas*able*.
```

Strikeout

Extension: strikeout

To strikeout a section of text with a horizontal line, begin and end it with ~~. Thus, for example,

```
This ~~is deleted text.~~
```

Superscripts and subscripts

Extension: superscript, subscript

Superscripts may be written by surrounding the superscripted text by ^ characters; subscripts may be written by surrounding the subscripted text by ~ characters. Thus, for example,

```
H~2~O is a liquid.  2^10^ is 1024.
```

If the superscripted or subscripted text contains spaces, these spaces must be escaped with backslashes. (This is to prevent accidental superscripting and subscripting through the ordinary use of ~ and ^.) Thus, if you want the letter P with 'a cat' in subscripts, use P~a\ cat~, not P~a cat~.

Verbatim

To make a short span of text verbatim, put it inside backticks:

```
What is the difference between `>>=` and `>>`?
```

If the verbatim text includes a backtick, use double backticks:

```
Here is a literal backtick `` ` ``.
```

(The spaces after the opening backticks and before the closing backticks will be ignored.)

The general rule is that a verbatim span starts with a string of consecutive backticks (optionally followed by a space) and ends with a string of the same number of backticks (optionally preceded by a space).

Note that backslash-escapes (and other markdown constructs) do not work in verbatim contexts:

```
This is a backslash followed by an asterisk: `\*`.
```

Extension: inline_code_attributes

Attributes can be attached to verbatim text, just as with fenced code blocks:

```
`<$>`{.haskell}
```

Small caps

To write small caps, you can use an HTML span tag:

```
<span style="font-variant:small-caps;">Small caps</span>
```

(The semicolon is optional and there may be space after the colon.) This will work in all output formats that support small caps.

Math

Extension: tex_math_dollars

Anything between two $ characters will be treated as TeX math. The opening $ must have a non-space character immediately to its right, while the closing $ must have a non-space character immediately to its left, and must not be followed immediately by a digit. Thus, \$20,000 and \$30,000 won't parse as math. If for some reason you need to enclose text in literal $ characters, backslash-escape them and they won't be treated as math delimiters.

TeX math will be printed in all output formats. How it is rendered depends on the output format:

Markdown, LaTeX, Org-Mode, ConTeXt It will appear verbatim between $ characters.

reStructuredText It will be rendered using an interpreted text role :math:, as described here

AsciiDoc It will be rendered as latexmath :[...] .

Texinfo It will be rendered inside a @math command.

groff man It will be rendered verbatim without $'s.

MediaWiki, DokuWiki It will be rendered inside <math> tags.

Textile It will be rendered inside <span class="math"> tags.

RTF, OpenDocument, ODT It will be rendered, if possible, using unicode characters, and will otherwise appear verbatim.

Docbook If the −−mathml flag is used, it will be rendered using mathml in an inlineequation or informalequation tag. Otherwise it will be rendered, if possible, using unicode characters.

Docx It will be rendered using OMML math markup.

FictionBook2 If the −−webtex option is used, formulas are rendered as images using Google Charts or other compatible web service, downloaded and embedded in the e-book. Otherwise, they will appear verbatim.

HTML, Slidy, DZSlides, S5, EPUB The way math is rendered in HTML will depend on the command-line options selected:

1. The default is to render TeX math as far as possible using unicode characters, as with RTF, DocBook, and OpenDocument output. Formulas are put inside a span with class="math", so that they may be styled differently from the surrounding text if needed.

2. If the −−latexmathml option is used, TeX math will be displayed between $ or $$ characters and put in <span> tags with class LaTeX. The LaTeXMathML script will be used to render it as formulas. (This trick does not work in all browsers, but it works in Firefox. In browsers that do not support LaTeXMathML, TeX math will appear verbatim between $ characters.)

3. If the −−jsmath option is used, TeX math will be put inside <span> tags (for inline math) or <div> tags (for display math) with class math. The jsMath script will be used to render it.

4. If the −−mimetex option is used, the mimeTeX CGI script will be called to generate images for each TeX formula. This should work in all browsers. The −−mimetex option takes an optional URL as argument. If no URL is specified, it will be assumed that the mimeTeX CGI script is at /cgi−bin/mimetex.cgi.

5. If the −−gladtex option is used, TeX formulas will be enclosed in <eq> tags in the HTML output. The resulting htex file may then be processed by gladTeX, which will produce image files for each formula and an html file with links to these images. So, the procedure is:

```
pandoc −s −−gladtex myfile.txt −o myfile.htex
gladtex −d myfile−images myfile.htex
# produces myfile.html and images in myfile−images
```

6. If the −−webtex option is used, TeX formulas will be converted to <img> tags that link to an external script that converts formulas to images. The formula will be URL-encoded and concatenated with the URL provided. If no URL is specified, the Google Chart API will be used (http://chart.apis.google.com/chart?cht=tx&chl=).

7. If the −−mathjax option is used, TeX math will be displayed between \(...\) (for inline math) or \[...\] (for display math) and put in <span> tags with class math. The MathJax script will be used to render it as formulas.

Raw HTML

Extension: raw_html

Markdown allows you to insert raw HTML (or DocBook) anywhere in a document (except verbatim contexts, where <, >, and & are interpreted literally). (Technically this is not an extension, since standard markdown allows it, but it has been made an extension so that it can be disabled if desired.)

The raw HTML is passed through unchanged in HTML, S5, Slidy, Slideous, DZSlides, EPUB, Markdown, and Textile output, and suppressed in other formats.

Extension: markdown_in_html_blocks

Standard markdown allows you to include HTML "blocks": blocks of HTML between balanced tags that are separated from the surrounding text with blank lines, and start and end at the left margin. Within these blocks, everything is interpreted as HTML, not markdown; so (for example), ∗ does not signify emphasis.

Pandoc behaves this way when the markdown_strict format is used; but by default, pandoc interprets material between HTML block tags as markdown. Thus, for example, Pandoc will turn

```
<table>
    <tr>
        <td>*one*</td>
        <td>[a link](http://google.com)</td>
    </tr>
</table>
```

into

```
<table>
    <tr>
        <td><em>one</em></td>
        <td><a href="http://google.com">a link</a></td>
    </tr>
</table>
```

whereas Markdown.pl will preserve it as is.

There is one exception to this rule: text between <script> and <style> tags is not interpreted as markdown.

This departure from standard markdown should make it easier to mix markdown with HTML block elements. For example, one can surround a block of markdown text with <div> tags without preventing it from being interpreted as markdown.

Extension: native_divs

Use native pandoc Div blocks for content inside <div> tags. For the most part this should give the same output as markdown_in_html_blocks, but it makes it easier to write pandoc filters to manipulate groups of blocks.

Extension: native_spans

Use native pandoc Span blocks for content inside <span> tags. For the most part this should give the same output as raw_html, but it makes it easier to write pandoc filters to manipulate groups of inlines.

Raw TeX

Extension: raw_tex

In addition to raw HTML, pandoc allows raw LaTeX, TeX, and ConTeXt to be included in a document. Inline TeX commands will be preserved and passed unchanged to the LaTeX and ConTeXt writers. Thus, for example, you can use LaTeX to include BibTeX citations:

```
This result was proved in \cite{jones.1967}.
```

Note that in LaTeX environments, like

```
\begin{tabular}{|l|l|}\hline
Age & Frequency \\ \hline
18--25  & 15 \\
26--35  & 33 \\
36--45  & 22 \\ \hline
\end{tabular}
```

the material between the begin and end tags will be interpreted as raw LaTeX, not as markdown.

Inline LaTeX is ignored in output formats other than Markdown, LaTeX, and ConTeXt.

LaTeX macros

Extension: latex_macros

For output formats other than LaTeX, pandoc will parse LaTeX \newcommand and \renewcommand definitions and apply the resulting macros to all LaTeX math. So, for example, the following will work in all output formats, not just LaTeX:

$\langle a, b, c \rangle$

In LaTeX output, the \newcommand definition will simply be passed unchanged to the output.

Links

Markdown allows links to be specified in several ways.

Automatic links

If you enclose a URL or email address in pointy brackets, it will become a link:

```
<http://google.com>
<sam@green.eggs.ham>
```

Inline links

An inline link consists of the link text in square brackets, followed by the URL in parentheses. (Optionally, the URL can be followed by a link title, in quotes.)

```
This is an [inline link](/url), and here's [one with
a title](http://fsf.org "click here for a good time!").
```

There can be no space between the bracketed part and the parenthesized part. The link text can contain formatting (such as emphasis), but the title cannot.

Email addresses in inline links are not autodetected, so they have to be prefixed with mailto:

```
[Write me!](mailto:sam@green.eggs.ham)
```

Reference links

An *explicit* reference link has two parts, the link itself and the link definition, which may occur elsewhere in the document (either before or after the link).

The link consists of link text in square brackets, followed by a label in square brackets. (There can be space between the two.) The link definition consists of the bracketed label, followed by a colon and a space, followed by the URL, and optionally (after a space) a link title either in quotes or in parentheses. The label must not be parseable as a citation (assuming the citations extension is enabled): citations take precedence over link labels.

Here are some examples:

```
[my label 1]: /foo/bar.html  "My title, optional"
[my label 2]: /foo
[my label 3]: http://fsf.org (The free software foundation)
[my label 4]: /bar#special  'A title in single quotes'
```

The URL may optionally be surrounded by angle brackets:

```
[my label 5]: <http://foo.bar.baz>
```

The title may go on the next line:

```
[my label 3]: http://fsf.org
  "The free software foundation"
```

Note that link labels are not case sensitive. So, this will work:

```
Here is [my link][FOO]

[Foo]: /bar/baz
```

In an *implicit* reference link, the second pair of brackets is empty, or omitted entirely:

```
See [my website][], or [my website].

[my website]: http://foo.bar.baz
```

Note: In Markdown.pl and most other markdown implementations, reference link definitions cannot occur in nested constructions such as list items or block quotes. Pandoc lifts this arbitrary seeming restriction. So the following is fine in pandoc, though not in most other implementations:

```
> My block [quote].
>
> [quote]: /foo
```

Internal links

To link to another section of the same document, use the automatically generated identifier (see Header identifiers in HTML, LaTeX, and ConTeXt, below). For example:

```
See the [Introduction](#introduction).
```

or

```
See the [Introduction].
```

```
[Introduction]: #introduction
```

Internal links are currently supported for HTML formats (including HTML slide shows and EPUB), LaTeX, and ConTeXt.

Images

A link immediately preceded by a ! will be treated as an image. The link text will be used as the image's alt text:

```
![la lune](lalune.jpg "Voyage to the moon")
```

```
![movie reel]
```

```
[movie reel]: movie.gif
```

Extension: implicit_figures

An image occurring by itself in a paragraph will be rendered as a figure with a caption.[4] (In LaTeX, a figure environment will be used; in HTML, the image will be placed in a div with class figure, together with a caption in a p with class caption.) The image's alt text will be used as the caption.

```
![This is the caption](/url/of/image.png)
```

If you just want a regular inline image, just make sure it is not the only thing in the paragraph. One way to do this is to insert a nonbreaking space after the image:

```
![This image won't be a figure](/url/of/image.png)\
```

Footnotes

Extension: footnotes

Pandoc's markdown allows footnotes, using the following syntax:

```
Here is a footnote reference,[^1] and another.[^longnote]
```

```
[^1]: Here is the footnote.
```

```
[^longnote]: Here's one with multiple blocks.

    Subsequent paragraphs are indented to show that they
belong to the previous footnote.

        { some.code }

    The whole paragraph can be indented, or just the first
    line.  In this way, multi-paragraph footnotes work like
    multi-paragraph list items.
```

---

[4]This feature is not yet implemented for RTF, OpenDocument, or ODT. In those formats, you'll just get an image in a paragraph by itself, with no caption.

```
This paragraph won't be part of the note, because it
isn't indented.
```

The identifiers in footnote references may not contain spaces, tabs, or newlines. These identifiers are used only to correlate the footnote reference with the note itself; in the output, footnotes will be numbered sequentially.

The footnotes themselves need not be placed at the end of the document. They may appear anywhere except inside other block elements (lists, block quotes, tables, etc.).

Extension: inline_notes

Inline footnotes are also allowed (though, unlike regular notes, they cannot contain multiple paragraphs). The syntax is as follows:

```
Here is an inline note.^[Inlines notes are easier to write, since
you don't have to pick an identifier and move down to type the
note.]
```

Inline and regular footnotes may be mixed freely.

Citations

Extension: citations

Using an external filter, pandoc−citeproc, pandoc can automatically generate citations and a bibliography in a number of styles. Basic usage is

```
pandoc −−filter pandoc−citeproc myinput.txt
```

In order to use this feature, you will need to specify a bibliography file using the bibliography metadata field in a YAML metadata section. The bibliography may have any of these formats:

Format	File extension
MODS	.mods
BibLaTeX	.bib
BibTeX	.bibtex
RIS	.ris
EndNote	.enl
EndNote XML	.xml
ISI	.wos
MEDLINE	.medline
Copac	.copac
JSON citeproc	.json

Note that .bib can generally be used with both BibTeX and BibLaTeX files, but you can use .bibtex to force BibTeX.

As an alternative to specifying a bibliography file, you can include the citation data directly in the references field of the document's YAML metadata. The field should contain an array of YAML-encoded references, for example:

```
−−−
references:
− id: fenner2012a
  title: One−click science marketing
```

```
  author:
- family: Fenner
  given: Martin
  container-title: Nature Materials
  volume: 11
  URL: 'http://dx.doi.org/10.1038/nmat3283'
  DOI: 10.1038/nmat3283
  issue: 4
  publisher: Nature Publishing Group
  page: 261-263
  type: article-journal
  issued:
    year: 2012
    month: 3
...
```

(The program mods2yaml, which comes with pandoc-citeproc, can help produce these from a MODS reference collection.)

By default, pandoc-citeproc will use a Chicago author-date format for citations and references. To use another style, you will need to specify a CSL 1.0 style file in the csl metadata field. A primer on creating and modifying CSL styles can be found at http://citationstyles.org/downloads/primer.html. A repository of CSL styles can be found at https://github.com/citation-style-language/styles. See also http://zotero.org/styles for easy browsing.

Citations go inside square brackets and are separated by semicolons. Each citation must have a key, composed of '@' + the citation identifier from the database, and may optionally have a prefix, a locator, and a suffix. The citation key must begin with a letter or _, and may contain alphanumerics, _, and internal punctuation characters (:.#$%&-+?<>~/). Here are some examples:

```
Blah blah [see @doe99, pp. 33-35; also @smith04, ch. 1].

Blah blah [@doe99, pp. 33-35, 38-39 and *passim*].

Blah blah [@smith04; @doe99].
```

A minus sign (-) before the @ will suppress mention of the author in the citation. This can be useful when the author is already mentioned in the text:

```
Smith says blah [-@smith04].
```

You can also write an in-text citation, as follows:

```
@smith04 says blah.

@smith04 [p. 33] says blah.
```

If the style calls for a list of works cited, it will be placed at the end of the document. Normally, you will want to end your document with an appropriate header:

```
last paragraph...
```

References

The bibliography will be inserted after this header. Note that the unnumbered class will be added to this header, so that the section will not be numbered.

If you want to include items in the bibliography without actually citing them in the body text, you can define a dummy nocite metadata field and put the citations there:

```
---
nocite: |
```

```
  @item1 ,  @item2
...
```

```
@item3
```

In this example, the document will contain a citation for item3 only, but the bibliography will contain entries for item1, item2, and item3.

Non-pandoc extensions

The following markdown syntax extensions are not enabled by default in pandoc, but may be enabled by adding +EXTENSION to the format name, where EXTENSION is the name of the extension. Thus, for example, markdown+hard_line_breaks is markdown with hard line breaks.

Extension: lists_without_preceding_blankline

Allow a list to occur right after a paragraph, with no intervening blank space.

Extension: hard_line_breaks

Causes all newlines within a paragraph to be interpreted as hard line breaks instead of spaces.

Extension: ignore_line_breaks

Causes newlines within a paragraph to be ignored, rather than being treated as spaces or as hard line breaks. This option is intended for use with East Asian languages where spaces are not used between words, but text is divided into lines for readability.

Extension: tex_math_single_backslash

Causes anything between \(and \) to be interpreted as inline TeX math, and anything between \[and \] to be interpreted as display TeX math. Note: a drawback of this extension is that it precludes escaping (and [.

Extension: tex_math_double_backslash

Causes anything between \\(and \\) to be interpreted as inline TeX math, and anything between \\[and \\] to be interpreted as display TeX math.

Extension: markdown_attribute

By default, pandoc interprets material inside block-level tags as markdown. This extension changes the behavior so that markdown is only parsed inside block-level tags if the tags have the attribute markdown=1.

Extension: mmd_title_block

Enables a MultiMarkdown style title block at the top of the document, for example:

```
Title :     My  t i t l e
Author :    John  Doe
Date :      September  1 ,  2008
Comment :  This  is  a  sample  mmd  t i t l e  block ,  with
          a  f i e l d  spanning  multiple  l i n e s .
```

See the MultiMarkdown documentation for details. If pandoc_title_block or yaml_metadata_block is enabled, it will take precedence over mmd_title_block.

Extension: abbreviations

Parses PHP Markdown Extra abbreviation keys, like

∗[HTML]: Hyper Text Markup Language

Note that the pandoc document model does not support abbreviations, so if this extension is enabled, abbreviation keys are simply skipped (as opposed to being parsed as paragraphs).

Extension: autolink_bare_uris

Makes all absolute URIs into links, even when not surrounded by pointy braces <...>.

Extension: ascii_identifiers

Causes the identifiers produced by auto_identifiers to be pure ASCII. Accents are stripped off of accented latin letters, and non-latin letters are omitted.

Extension: link_attributes

Parses multimarkdown style key-value attributes on link and image references. Note that pandoc's internal document model provides nowhere to put these, so they are presently just ignored.

Extension: mmd_header_identifiers

Parses multimarkdown style header identifiers (in square brackets, after the header but before any trailing #s in an ATX header).

Extension: compact_definition_lists

Activates the definition list syntax of pandoc 1.12.x and earlier. This syntax differs from the one described above in several respects:

- No blank line is required between consecutive items of the definition list.
- To get a "tight" or "compact" list, omit space between consecutive items; the space between a term and its definition does not affect anything.
- Lazy wrapping of paragraphs is not allowed: the entire definition must be indented four spaces.[5]

Markdown variants

In addition to pandoc's extended markdown, the following markdown variants are supported:

> footnotes, pipe_tables, raw_html, markdown_attribute, fenced_code_blocks, definition_lists , intraword_underscores, header_attributes, abbreviations.

---

[5]To see why laziness is incompatible with relaxing the requirement of a blank line between items, consider the following example: bar : definition foo : definition Is this a single list item with two definitions of "bar," the first of which is lazily wrapped, or two list items? To remove the ambiguity we must either disallow lazy wrapping or require a blank line between list items.

pipe_tables, raw_html, tex_math_single_backslash, fenced_code_blocks, auto_identifiers , ascii_identifiers , backtick_code_blocks, autolink_bare_uris, intraword_underscores, strikeout , hard_line_breaks

pipe_tables raw_html, markdown_attribute, link_attributes, raw_tex, tex_math_double_backslash, intraword_underscores, mmd_title_block, footnotes, definition_lists , all_symbols_escapable, implicit_header_reference auto_identifiers , mmd_header_identifiers

raw_html

Extensions with formats other than markdown

Some of the extensions discussed above can be used with formats other than markdown:

markdown_p(in(Multhitab)Maakhtidfiyon)atridtoboroghithmutex, rst, mediawiki, and textile input (and is used by default).

- tex_math_dollars, tex_math_single_backslash, and tex_math_double_backslash can be used with html input. (This is handy for reading web pages formatted using MathJax, for example.)

Chapter 6

Producing slide shows with Pandoc

You can use Pandoc to produce an HTML + javascript slide presentation that can be viewed via a web browser. There are five ways to do this, using S5, DZSlides, Slidy, Slideous, or reveal.js. You can also produce a PDF slide show using LaTeX beamer.

Here's the markdown source for a simple slide show, habits.txt:

```
% Habits
% John Doe
% March 22, 2005

# In the morning

## Getting up

- Turn off alarm
- Get out of bed

## Breakfast

- Eat eggs
- Drink coffee

# In the evening

## Dinner

- Eat spaghetti
- Drink wine

------------------

![picture of spaghetti](images/spaghetti.jpg)

## Going to sleep

- Get in bed
- Count sheep
```

To produce an HTML/javascript slide show, simply type

```
pandoc -t FORMAT -s habits.txt -o habits.html
```

where FORMAT is either s5, slidy, slideous, dzslides, or revealjs.

For Slidy, Slideous, reveal.js, and S5, the file produced by pandoc with the −s/−−standalone option embeds a link to javascripts and CSS files, which are assumed to be available at the relative path s5/default (for S5), slideous (for Slideous), reveal.js (for reveal.js), or at the Slidy website at w3.org (for Slidy). (These paths can be changed by setting the slidy−url, slideous−url, revealjs−url, or s5−url variables; see −−variable, above.) For DZSlides, the (relatively short) javascript and css are included in the file by default.

With all HTML slide formats, the −−self−contained option can be used to produce a single file that contains all of the data necessary to display the slide show, including linked scripts, stylesheets, images, and videos.

To produce a PDF slide show using beamer, type

```
pandoc −t beamer habits.txt −o habits.pdf
```

Note that a reveal.js slide show can also be converted to a PDF by printing it to a file from the browser.

Structuring the slide show

By default, the *slide level* is the highest header level in the hierarchy that is followed immediately by content, and not another header, somewhere in the document. In the example above, level 1 headers are always followed by level 2 headers, which are followed by content, so 2 is the slide level. This default can be overridden using the −−slide−level option.

The document is carved up into slides according to the following rules:

- A horizontal rule always starts a new slide.

- A header at the slide level always starts a new slide.

- Headers *below* the slide level in the hierarchy create headers *within* a slide.

- Headers *above* the slide level in the hierarchy create "title slides," which just contain the section title and help to break the slide show into sections.

- A title page is constructed automatically from the document's title block, if present. (In the case of beamer, this can be disabled by commenting out some lines in the default template.)

These rules are designed to support many different styles of slide show. If you don't care about structuring your slides into sections and subsections, you can just use level 1 headers for all each slide. (In that case, level 1 will be the slide level.) But you can also structure the slide show into sections, as in the example above.

Note: in reveal.js slide shows, if slide level is 2, a two-dimensional layout will be produced, with level 1 headers building horizontally and level 2 headers building vertically. It is not recommended that you use deeper nesting of section levels with reveal.js.

Incremental lists

By default, these writers produces lists that display "all at once." If you want your lists to display incrementally (one item at a time), use the −i option. If you want a particular list to depart from the default (that is, to display incrementally without the −i option and all at once with the −i option), put it in a block quote:

```
> − Eat spaghetti
> − Drink wine
```

In this way incremental and nonincremental lists can be mixed in a single document.

Inserting pauses

You can add "pauses" within a slide by including a paragraph containing three dots, separated by spaces:

```
# Slide with a pause

content before the pause

. . .

content after the pause
```

Styling the slides

You can change the style of HTML slides by putting customized CSS files in $DATADIR/s5/default (for S5), $DATADIR/slidy (for Slidy), or $DATADIR/slideous (for Slideous), where $DATADIR is the user data directory (see −−data−dir, above). The originals may be found in pandoc's system data directory (generally $CABALDIR/pandoc−VERSION/s5/default). Pandoc will look there for any files it does not find in the user data directory.

For dzslides, the CSS is included in the HTML file itself, and may be modified there.

For reveal.js, themes can be used by setting the theme variable, for example:

```
−V theme=moon
```

Or you can specify a custom stylesheet using the −−css option.

To style beamer slides, you can specify a beamer "theme" or "colortheme" using the −V option:

```
pandoc −t beamer habits.txt −V theme:Warsaw −o habits.pdf
```

Note that header attributes will turn into slide attributes (on a <div> or <section>) in HTML slide formats, allowing you to style individual slides. In Beamer, the only header attribute that affects slides is the allowframebreaks class, which sets the allowframebreaks option, causing multiple slides to be created if the content overfills the frame. This is recommended especially for bibliographies:

```
# References {.allowframebreaks}
```

Speaker notes

reveal.js has good support for speaker notes. You can add notes to your markdown document thus:

```
<div class="notes">
This is my note.

− It can contain markdown
− like this list

</div>
```

To show the notes window, press s while viewing the presentation. Notes are not yet supported for other slide formats, but the notes will not appear on the slides themselves.

Chapter 7

EPUB Metadata

EPUB metadata may be specified using the −−epub−metadata option, but if the source document is markdown, it is better to use a YAML metadata block. Here is an example:

```
−−−
title :
− type : main
  text : My Book
− type : subtitle
  text : An investigation of metadata
creator :
− role : author
  text : John Smith
− role : editor
  text : Sarah Jones
identifier :
− scheme : DOI
  text : doi:10.234234.234/33
publisher : My Press
rights : (C) 2007 John Smith , CC BY–NC
...
```

The following fields are recognized:

Either a string value or an object with fields text and scheme. Valid values for scheme are ISBN−10, GTIN−13, UPC, ISMN−10, DOI, LCCN, GTIN−14, ISBN−13, Legal deposit number, URN, OCLC, ISMN−13, ISBN−A, JP, OLCC.

Either a string value, or an object with fields file −as and type, or a list of such objects. Valid values for type are main, subtitle , short, collection , edition, extended.

Either a string value, or an object with fields role, file −as, and text, or a list of such objects. Valid values for role are marc relators, but pandoc will attempt to translate the human-readable versions (like "author" and "editor") to the appropriate marc relators.

Same format as creator.

A string value in YYYY−MM−DD format. (Only the year is necessary.) Pandoc will attempt to convert other common date formats.

A string value in RFC5646 format. Pandoc will default to the local language if nothing is specified.

A string value or a list of such values.

A string value.

A string value.

A string value.

A string value.

A string value.

A string value.

A string value (path to cover image).

A string value (path to CSS stylesheet).

Either ltr or rtl. Specifies the page−progression−direction spine attribute.

Chapter 8

Literate Haskell support

If you append +lhs (or +literate_haskell) to an appropriate input or output format (markdown, markdown_strict, rst, or latex for input or output; beamer, html or html5 for output only), pandoc will treat the document as literate Haskell source. This means that

- In markdown and Markdown input, "bird track" sections will be parsed as Haskell code rather than block quotations. Text between \begin{code} and \end{code} will also be treated as Haskell code.

- In markdown output, code blocks with classes haskell and literate will be rendered using bird tracks, and block quotations will be indented one space, so they will not be treated as Haskell code. In addition, headers will be rendered setext-style (with underlines) rather than atx-style (with '#' characters). (This is because ghc treats '#' characters in column 1 as introducing line numbers.)

- In restructured text input, "bird track" sections will be parsed as Haskell code.

- In restructured text output, code blocks with class haskell will be rendered using bird tracks.

- In LaTeX input, text in code environments will be parsed as Haskell code.

- In LaTeX output, code blocks with class haskell will be rendered inside code environments.

- In HTML output, code blocks with class haskell will be rendered with class literatehaskell and bird tracks.

Examples:

```
pandoc −f markdown+lhs −t html
```

reads literate Haskell source formatted with markdown conventions and writes ordinary HTML (without bird tracks).

```
pandoc −f markdown+lhs −t html+lhs
```

writes HTML with the Haskell code in bird tracks, so it can be copied and pasted as literate Haskell source.

Chapter 9

Custom writers

Pandoc can be extended with custom writers written in lua. (Pandoc includes a lua interpreter, so lua need not be installed separately.)

To use a custom writer, simply specify the path to the lua script in place of the output format. For example:

```
pandoc −t data/sample.lua
```

Creating a custom writer requires writing a lua function for each possible element in a pandoc document. To get a documented example which you can modify according to your needs, do

```
pandoc −−print−default−data−file sample.lua
```

Chapter 10

Authors

© 2006-2013 John MacFarlane (jgm at berkeley dot edu). Released under the GPL, version 2 or greater. This software carries no warranty of any kind. (See COPYRIGHT for full copyright and warranty notices.) Other contributors include Recai Oktaş, Paulo Tanimoto, Peter Wang, Andrea Rossato, Eric Kow, infinity0x, Luke Plant, shreevatsa.public, Puneeth Chaganti, Paul Rivier, rodja.trappe, Bradley Kuhn, thsutton, Nathan Gass, Jonathan Daugherty, Jérémy Bobbio, Justin Bogner, qerub, Christopher Sawicki, Kelsey Hightower, Masayoshi Takahashi, Antoine Latter, Ralf Stephan, Eric Seidel, B. Scott Michel, Gavin Beatty, Sergey Astanin, Arlo O'Keeffe, Denis Laxalde, Brent Yorgey, David Lazar, Jamie F. Olson, Matthew Pickering, Albert Krewinkel, mb21, Jesse Rosenthal.